IMU-Stabilized Wide-Angle Star Tracker for Autonomous Navigation in Light-Polluted Environments

Technical Review and System Architecture

Giray Yıllıkçı

STP606E Term Project

Abstract

This paper presents a novel approach to star-based autonomous navigation for unmanned vehicles operating in challenging environments characterized by significant light pollution, platform vibration, and dynamic motion. Traditional star trackers require dark skies and stable platforms, limiting their application in sea-level operations and mobile platforms. We propose an integrated system combining ultra-wide-angle optics (140-220 degrees field of view), tactical-grade inertial measurement units (IMU), and advanced image processing to enable celestial navigation under adverse conditions. The system employs short exposure times (3-5 seconds) synchronized with high-frequency IMU data acquisition to facilitate post-processing stabilization and star detection. We analyze two fundamental approaches: video frame stacking with IMU-based alignment versus direct long-exposure imaging with computational motion compensation. The system architecture addresses key challenges including wide-angle lens distortion, light pollution filtering, vibration-induced motion blur, and real-time attitude determination. Performance analysis demonstrates feasibility for autonomous navigation applications across unmanned aerial vehicles (UAVs), unmanned ground vehicles (UGVs), and unmanned surface vehicles (USVs). This work contributes to expanding the operational envelope of celestial navigation systems beyond traditional aerospace applications.

1. Introduction

Celestial navigation has served as a cornerstone of positioning and orientation determination for centuries, from maritime navigation to modern spacecraft attitude control. However, traditional star tracker systems operate under assumptions that limit their applicability to mobile autonomous platforms: dark sky conditions (typically magnitude 6+ stars visible), stable mounting platforms with minimal vibration, and relatively narrow fields of view optimized for specific star catalog matching algorithms.

The proliferation of unmanned autonomous systems across maritime, terrestrial, and aerial domains has created demand for navigation solutions that combine the independence of celestial navigation with robustness to GPS denial, jamming, or spoofing. While traditional star trackers excel in aerospace applications, they face significant challenges when adapted to sea-level operations where light pollution from urban areas, industrial facilities, and vehicle lighting can wash out all but the brightest celestial objects.

This paper addresses the fundamental question: can star-based navigation be made practical for autonomous vehicles operating at sea level, in light-polluted environments, while subject to continuous vibration and dynamic motion? We propose a system architecture that deliberately trades traditional star tracker design principles for approaches better suited to these challenging conditions. Specifically, we employ ultra-wide-angle optics to maximize sky coverage and stellar observations per frame, accept short integration times that individually fail to resolve faint stars, and rely on synchronized IMU data to enable computational stabilization and frame stacking.

1.1 Motivation and Applications

Modern autonomous systems face an increasing threat landscape regarding position, navigation, and timing (PNT) services. GPS/GNSS systems, while ubiquitous, present single points of failure vulnerable to intentional jamming, spoofing, or natural interference. Military operations increasingly assume contested electromagnetic environments where satellite navigation may be denied or unreliable. Commercial autonomous vehicles, from delivery drones to autonomous ships, require redundant navigation sources to ensure safety and operational continuity.

Celestial navigation offers inherent advantages for these scenarios. Stars cannot be jammed, spoofed, or denied by adversarial action. The star field provides a globally available, passive reference frame requiring no external infrastructure. Unlike vision-based simultaneous localization and mapping (SLAM) systems that require prior maps or distinctive landmarks, celestial navigation works anywhere on Earth with sky visibility. However, realizing these advantages in practical autonomous systems requires overcoming the environmental and platform limitations that constrain conventional star tracker designs.

Primary applications for this technology include autonomous maritime vessels operating in coastal waters where light pollution is significant, ground robots navigating in urban or suburban environments, and UAVs flying at low altitudes where atmospheric scattering and ground-based lighting reduce star visibility. The system must function as a complementary navigation source, providing attitude determination and supporting position estimation when integrated with other sensors.

1.2 Key Technical Challenges

Adapting star tracking to mobile platforms in light-polluted environments presents several interconnected challenges. Light pollution fundamentally reduces the limiting magnitude of visible stars, potentially from magnitude 6 in dark sites to magnitude 2-3 in urban environments. This reduction dramatically decreases the number of observable stars, from thousands to dozens, complicating star pattern recognition and attitude determination algorithms.

Platform vibration introduces motion blur that can exceed the angular size of point sources during exposure. A vehicle traveling over rough terrain or a vessel in moderate seas may experience vibrations at frequencies from 1-100 Hz with angular amplitudes from milliradians to degrees. Traditional solutions involve vibration isolation, which adds mass, complexity, and cost. Our approach instead embraces the vibration, capturing it with high-rate IMU sampling and compensating computationally.

Wide-angle optics necessary for maximizing stellar observations introduce severe barrel distortion and vignetting. Lenses approaching 180-degree field of view exhibit radial distortion that can displace star positions by tens of pixels from their rectified locations. This distortion must be accurately characterized and compensated to enable star catalog matching. Additionally, wide-angle lenses generally have lower f-numbers at field edges, reducing light collection efficiency precisely where observations are most valuable.

The core trade-off in exposure strategy centers on motion blur versus signal accumulation. Shorter exposures freeze platform motion but collect insufficient light to detect faint stars in individual frames. Longer exposures accumulate light but allow platform motion to blur point sources into streaks. Our proposed approach of 3-5 second exposures with IMU-synchronized data acquisition attempts to balance these constraints, accepting that individual frames will be noisy while enabling post-processing recovery through temporal integration.

2. Background and Related Work

2.1 Traditional Star Tracker Systems

Spacecraft star trackers represent the most mature application of celestial attitude determination. These systems typically employ narrow to moderate field-of-view optics (5-30 degrees), medium resolution sensors (1-2 megapixels), and integration times from milliseconds to seconds depending on spacecraft dynamics. Star identification algorithms such as grid, pyramid, and geometric voting methods achieve high reliability by matching observed star patterns against onboard catalogs.

Spacecraft benefit from ideal operating conditions: zero atmospheric extinction, minimal scattered light (except near bright bodies), and generally low angular rates outside maneuvers. Vibration is largely absent, allowing sharp stellar images. These systems achieve attitude determination accuracy of arcseconds to arcminutes, with update rates of 1-10 Hz. However, their design assumptions break down completely for ground-based or low-altitude aerial applications.

Ground-based astronomical applications have developed extensive techniques for dealing with atmospheric effects, light pollution, and tracking celestial objects. Adaptive optics, lucky imaging, and speckle interferometry address atmospheric turbulence. Long-exposure tracking mounts compensate for Earth's rotation. However, these solutions assume stationary installations and cannot transfer directly to mobile platforms.

2.2 IMU-Aided Vision Systems

The integration of inertial measurement units with optical sensors has proven valuable across numerous applications. Visual-inertial odometry (VIO) systems combine IMU measurements with feature tracking to estimate ego-motion in GPS-denied environments. The IMU provides high-rate motion estimates between image frames, enabling prediction of feature locations and handling rapid movements that would otherwise cause tracking failure.

IMU quality significantly impacts system performance. Consumer-grade MEMS IMUs exhibit bias drift of 10-100 degrees per hour, requiring frequent visual corrections. Tactical-grade IMUs such as those in the Orange Cube platform achieve bias stability of 1-10 degrees per hour with lower noise density, enabling longer intervals between visual updates. For our application, tactical-grade performance provides sufficient fidelity to track platform orientation during 3-5 second exposures while maintaining sub-degree accuracy.

Image stabilization in consumer cameras increasingly relies on IMU data. Both optical image stabilization (OIS) and electronic image stabilization (EIS) use gyroscope measurements to compensate for hand shake. These systems demonstrate that IMU-based stabilization can work effectively for exposure times up to several seconds. However, they typically address high-frequency vibration (5-20 Hz hand tremor) rather than the broader frequency range experienced by mobile vehicles.

2.3 Light Pollution Mitigation

Light pollution presents both a radiometric challenge (reduced signal-to-noise ratio) and a spatial challenge (gradients in sky background brightness). Urban skyglow results primarily from sodium vapor lamps (589 nm), metal halide lamps (broad spectrum), and increasingly LED lighting (blue-rich spectrum). The spectrum of light pollution enables some mitigation through narrowband filtering, though this approach sacrifices overall light collection.

Computational approaches to light pollution compensation include background subtraction, high-pass filtering, and adaptive thresholding. Dark frame subtraction removes sensor dark current and fixed pattern noise. Flat field correction compensates for vignetting and sensor response variations. However, variable sky background from light pollution domes requires dynamic background modeling, potentially on a per-frame basis for moving platforms.

Recent work in urban astronomy has demonstrated that even in heavily light-polluted environments (Bortle class 8-9), stars brighter than magnitude 3 remain detectable with appropriate imaging techniques. This limiting magnitude, while drastically reduced from dark sky conditions, still provides access to several hundred stars across the full sky, sufficient for attitude determination if properly distributed across the field of view.

3. System Architecture

3.1 Hardware Configuration

The proposed system comprises three primary hardware components: optical subsystem, IMU, and processing platform. The optical subsystem employs a fisheye or ultra-wide-angle lens with 140-220 degree field of view paired with a sensitive CMOS or CCD sensor. Wide aperture (f/2.8 or faster) maximizes light collection despite the challenging light pollution environment. Sensor resolution should balance spatial resolution with pixel size, typically 2-4 megapixels with 3-5 micron pixels to maintain adequate well depth and low read noise.

The IMU selection critically impacts system performance. Tactical-grade MEMS IMUs such as the InvenSense ICM-42688 or comparable sensors provide gyroscope noise density below 0.004 degrees per second per root-Hertz with bias instability under 2 degrees per hour. Accelerometer performance, while less critical for attitude estimation, aids in detecting certain motion artifacts and supports integration with other navigation sensors. The IMU must operate at rates of 100-400 Hz to capture platform dynamics with sufficient fidelity for post-processing compensation.

Time synchronization between camera and IMU requires precision below 10 milliseconds to prevent attitude errors that would misalign star positions during frame stacking. Hardware synchronization through trigger signals ensures deterministic timing. Each image frame captures a precise timestamp, with IMU data buffered at high rate around exposure intervals. This architecture enables precise temporal alignment in post-processing.

3.2 Operational Modes

The system supports two primary operational modes based on processing approach. In video stacking mode, the camera captures continuous video at 15-30 frames per second with exposure times of 100-200 milliseconds per frame. IMU data streams continuously at 200 Hz. Post-processing aligns and stacks frames using IMU-derived rotation estimates, effectively increasing integration time while maintaining alignment to a reference frame. This approach provides flexibility in final integration time and enables rejection of frames affected by transient interference such as aircraft lights or lightning.

In long-exposure mode, the camera captures individual frames with 3-5 second exposures. IMU data still streams at 200 Hz throughout exposure. Post-processing uses the IMU trajectory to deblur the resulting image, effectively reversing the motion blur introduced by platform dynamics. This mode reduces data volume and processing load but offers less flexibility in handling transient interference. The fundamental trade-off between these modes centers on whether discrete frame alignment or continuous motion deblurring provides superior performance for a given operational scenario.

3.3 Processing Pipeline

The processing pipeline transforms raw image data and IMU measurements into attitude estimates through several stages. Initial preprocessing applies dark frame subtraction, flat field correction, and lens distortion compensation. The distortion model, calibrated offline using stellar references or checkerboard targets, maps observed pixel coordinates to rectified angular coordinates. For fisheye lenses, equidistant or stereographic projection models typically provide adequate fit.

IMU processing integrates gyroscope measurements to estimate platform orientation throughout each exposure or video sequence. A quaternion-based integration scheme avoids gimbal lock while maintaining computational efficiency. Gyroscope bias estimation can employ several approaches: pre-integration calibration, estimation during platform stationary periods, or continuous estimation using detected stars as orientation references. The estimated orientation trajectory feeds into the image alignment or deblurring stage.

For video stacking mode, each frame undergoes rotation to a common reference frame using the IMU-estimated attitude difference. Bilinear or bicubic interpolation resamples pixel values. Accumulated frames combine through median stacking to reject outliers or weighted averaging to minimize noise. The resulting synthetic long-exposure image proceeds to star detection. For long-exposure mode, deconvolution using the IMU-derived point spread function sharpens star images. Lucy-Richardson or Wiener deconvolution algorithms can recover point sources from motion-blurred streaks, though computational cost increases with trajectory complexity.

Star detection employs adaptive thresholding to handle spatially varying background brightness from light pollution. Local background estimation in windows of 50-100 pixels enables detection of stars against bright regions of sky. Centroiding of detected sources using center-of-mass or Gaussian fitting provides sub-pixel position accuracy. Detected stars then undergo pattern matching against star catalogs to establish correspondence between image features and known celestial objects.

4. IMU-Based Image Stabilization Analysis

4.1 Gyroscope Integration and Attitude Estimation

Accurate attitude estimation from gyroscope measurements forms the foundation of image stabilization. The attitude quaternion evolves according to the kinematic differential equation relating angular velocity to quaternion derivative. For discrete-time implementation at sampling interval Δt, first-order integration provides adequate accuracy for $\Delta t < 5$ milliseconds. Higher-order integration schemes such as fourth-order Runge-Kutta improve accuracy for larger time steps but increase computational cost.

Gyroscope measurements contain several error sources that impact attitude estimation accuracy. White noise, characterized by angle random walk (ARW) coefficient, accumulates as the square root of integration time. For tactical-grade gyroscopes with ARW of 0.1 degrees per root-hour, attitude uncertainty from noise reaches 0.01 degrees over a 3-second exposure. Bias instability, typically specified at 2-10 degrees per hour for tactical-grade units, contributes linear attitude error over time. Temperature-induced bias variations can reach 0.1-1 degrees per degree Celsius, necessitating thermal compensation or controlled environments.

Scale factor errors and axis misalignment introduce attitude errors proportional to rotation magnitude. For gyroscopes specified at 0.1 percent scale factor error, a 10-degree rotation accumulates 0.01-degree attitude error. Cross-axis sensitivity can reach 1-2 percent, coupling rotation about one axis into measurements of perpendicular axes. Calibration procedures can characterize and partially compensate these deterministic errors, though residual errors remain after calibration.

4.2 Frame Alignment Techniques

Video stacking mode requires aligning each frame to a common reference frame using IMU-estimated rotation. The alignment process must account for lens distortion, performing rotation in rectified angular space rather than pixel space to avoid coupling distortion with rotation. For each pixel in the output aligned frame, the algorithm computes the corresponding viewing direction, applies the rotation inverse, accounts for lens distortion, and samples the input frame at the resulting pixel location.

Computational efficiency becomes critical for real-time or near-real-time operation. For a 2-megapixel sensor at 30 frames per second, naive implementation requires 60 million pixel rotations per second. Optimization strategies include lookup table generation for distortion correction, GPU acceleration of rotation and resampling operations, and hierarchical processing where initial coarse alignment precedes fine alignment of regions containing potential stars.

Cumulative alignment error results from both IMU drift and resampling artifacts. IMU drift primarily affects absolute orientation estimation rather than relative frame-to-frame alignment. Over a 10-second video sequence, tactical-grade gyroscope drift remains below 0.03 degrees, small compared to pixel resolution for wide-angle optics. Resampling introduces high-frequency noise and can slightly blur features, particularly when interpolating between high-contrast pixels. Median stacking of aligned frames provides some rejection of resampling artifacts.

4.3 Motion Deblurring for Long Exposures

Long-exposure imaging with motion deblurring presents an alternative to frame stacking. During a 3-5 second exposure with platform vibration, each star traces a path across the sensor, creating a streak rather than a point. The IMU trajectory during exposure defines the point spread function (PSF) that blurred the ideal point source into the observed streak. Deconvolution algorithms attempt to invert this blurring process, recovering the underlying point source distribution.

The deblurring problem becomes well-posed when the PSF is known with sufficient accuracy. IMU-derived PSF depends on gyroscope measurement accuracy and integration fidelity. PSF errors from IMU measurement noise and bias manifest as residual blur in deblurred images. For a 3-second exposure with 0.01-degree RMS attitude error, residual blur spans approximately 2-3 pixels for a wide-angle lens, degrading but not preventing star detection.

Wiener deconvolution provides a straightforward approach, balancing deblurring against noise amplification through a regularization parameter. However, it assumes spatially uniform PSF, which breaks down for wide-angle lenses where stars at different field positions experience different motion trajectories. Spatially variant deblurring divides the image into regions, computing local PSF for each region based on the viewing direction and platform motion. This approach increases computational cost but significantly improves results for wide field-of-view systems.

Lucy-Richardson deconvolution, an iterative maximum-likelihood approach, often outperforms Wiener deconvolution for point sources like stars. It naturally enforces non-negativity and can incorporate Poisson noise models appropriate for photon counting. Typical convergence requires 10-50 iterations, with computation time proportional to image size and PSF extent. For our application, PSF extent depends on platform vibration magnitude, potentially spanning 10-100 pixels for severe motion.

5. Video Stacking vs Long Exposure: Comparative Analysis

5.1 Signal-to-Noise Ratio Considerations

Signal-to-noise ratio (SNR) fundamentally determines star detection capability. In video stacking mode with N frames of individual exposure time t, total integration time equals Nt. Read noise contributes independently to each frame, accumulating as sqrt(N) times single-frame read noise. For typical CMOS sensors with 3-5 electron read noise, 30 frames contribute sqrt(30) ≈ 5.5 times read noise, or approximately 17-27 electrons total. Photon shot noise, following Poisson statistics, accumulates as sqrt(Nt × photon rate), identical to a single long exposure.

Long-exposure mode incurs read noise only once per integration period, providing an SNR advantage when read noise dominates. For a 3-second exposure equivalent to 30 frames at 100ms each, read noise is 5.5 times lower in long-exposure mode. However, this advantage diminishes for sensors with low read noise or when sky background (light pollution) dominates noise. In light-polluted environments, sky background generates hundreds to thousands of electrons per second per pixel, dwarfing read noise contribution.

Dark current, while less significant for cooled sensors, contributes thermal noise proportional to total exposure time. Both modes accumulate identical dark current over equivalent integration times. Temperature-stabilized sensors at 0-10 degrees Celsius exhibit dark current of 0.01-0.1 electrons per pixel per second, contributing negligibly compared to sky background in light-polluted conditions.

5.2 Motion Blur and Temporal Resolution

Video stacking mode with short individual exposures provides natural motion blur suppression within each frame. A 100ms exposure with platform vibration frequency of 10 Hz captures approximately one oscillation cycle, limiting motion blur to the vibration amplitude. For 0.1-degree vibration amplitude, blur spans roughly 5-10 pixels for a 180-degree FOV fisheye lens on a 2-megapixel sensor. This limited blur enables direct star detection in individual frames before stacking, facilitating outlier rejection.

Long-exposure mode accumulates motion blur throughout the exposure period. A 3-second exposure captures 30 vibration cycles at 10 Hz, creating complex streak patterns. While IMU data enables deblurring, the process faces challenges from non-uniform motion, IMU errors, and computational cost. Deblurring algorithms struggle with highly curved trajectories resulting from multi-frequency vibration components. Success depends critically on IMU accuracy relative to the blur extent.

Temporal resolution differs significantly between modes. Video mode naturally segments time into discrete frames, enabling rejection of frames containing transient interference like aircraft lights, lightning, or laser illumination. Each frame provides an independent observation, facilitating statistical outlier detection. Long-exposure mode accumulates all photons during exposure, with no ability to exclude transient events post-capture. In environments with frequent transient interference, this distinction strongly favors video stacking.

5.3 Processing Complexity and Latency

Computational requirements differ substantially between approaches. Video stacking processes N frames, each requiring distortion correction, rotation, resampling, and stacking. For 30 frames of 2-megapixel resolution, this totals 60 megapixel operations plus stacking. Modern GPUs handle this workload in seconds, enabling near-real-time processing. Memory requirements scale with N and frame size, reaching several hundred megabytes for typical scenarios.

Long-exposure mode performs single-frame distortion correction followed by deblurring. Deblurring proves computationally expensive, with cost scaling as $O(MN \log(MN))$ for FFT-based methods where $M \times N$ represents image dimensions. Spatially variant deblurring multiplies this cost by the number of regions. For a 2-megapixel image divided into 100 regions, processing can require tens of seconds on CPU or several seconds on GPU. However, total data volume is lower, and storage requirements decrease proportionally to the number of frames avoided.

Latency considerations favor long-exposure mode when pipeline latency matters. Video stacking cannot begin until all frames are captured, introducing latency equal to the total integration time. Long-exposure mode completes capture in one exposure period, with processing beginning immediately after. For a 3-second integration requirement, video mode at 30 fps introduces 3 seconds of capture latency, while long-exposure mode enables processing to start 3 seconds earlier, partially offsetting its higher computational cost.

5.4 Robustness and Failure Modes

Failure mode analysis reveals complementary strengths. Video stacking degrades gracefully with IMU errors. Small attitude errors misalign frames by corresponding pixel amounts, slightly blurring the stacked result. Star detection remains possible even with 2-3 pixel misalignment, though centroiding accuracy suffers. Complete IMU failure prevents frame alignment but individual frames may still contain detectable stars, enabling fallback to single-frame processing.

Long-exposure mode exhibits more brittle failure characteristics. IMU errors directly corrupt the deblurring PSF, potentially causing deconvolution to diverge or amplify noise rather than recovering point sources. Severe IMU drift can render deblurring completely ineffective, leaving only motion-blurred streaks. Recovery strategies include attempting star detection directly on streaks using Radon transform methods, though this approach struggles with overlapping streaks and provides reduced attitude accuracy.

Both modes face challenges from saturation. Bright stars or planets can saturate sensor pixels, creating blooming artifacts that corrupt adjacent pixels. Video mode can detect and exclude saturated frames, while long-exposure mode must handle saturation through pre-exposure checks or multi-exposure HDR techniques. Light pollution variations across the field of view can cause some regions to approach saturation while others remain underexposed, complicating exposure selection for long-exposure mode.

6. Star Detection and Pattern Matching Algorithms

6.1 Adaptive Star Detection in Light-Polluted Images

Star detection in light-polluted environments requires adaptive approaches that account for spatially varying background brightness. Global thresholding fails when sky background varies by factors of 10-100 across the field of view, as commonly occurs near horizon or when artificial light sources illuminate portions of the scene. Local background estimation divides the image into overlapping regions, computing median or low-percentile background level for each region. Threshold for star detection then adapts to local background plus a multiple of estimated noise.

Morphological operations can separate point sources from extended background. Top-hat filtering, applying morphological opening followed by subtraction from the original image, removes low-frequency background while preserving point sources. This approach proves effective even when background varies smoothly across the field. Matched filtering with a Gaussian kernel approximating the point spread function provides optimal detection in Gaussian noise, though computational cost increases with kernel size.

False positive rejection becomes critical in light-polluted environments where noise fluctuations and artifacts can mimic faint stars. Size-based filtering eliminates detections much larger or smaller than expected stellar point spread functions. Elongation testing removes streak artifacts from aircraft or satellites. Statistical hypothesis testing using likelihood ratios can distinguish genuine stars from noise peaks, though this requires accurate noise models accounting for Poisson statistics, read noise, and quantization.

6.2 Star Pattern Recognition for Wide-Angle Optics

Traditional star pattern recognition algorithms such as triangle matching and planar angle methods assume moderate field-of-view optics with negligible distortion. Ultra-wide-angle lenses violate these assumptions, requiring adapted approaches. Geometric hashing methods can accommodate distortion by working in rectified angular coordinates rather than pixel space. Each detected star maps to a unit vector in the camera reference frame after distortion correction, with angular separation between stars providing distortion-invariant features.

Pyramid algorithms examine sets of four stars, computing ratios of inter-star angular separations to create rotation-invariant descriptors. Hash tables index catalog stars by these descriptors, enabling rapid lookup of candidate matches. For a 180-degree field of view potentially containing 50-200 stars in light-polluted conditions, the combinatorial space of four-star patterns reaches millions. Efficient implementation requires careful pruning of unlikely combinations based on brightness ordering and spatial distribution.

Lost-in-space algorithms must identify star patterns without prior attitude knowledge. This scenario occurs during system initialization or after extended periods without valid solutions. Robustness requires successful matching with as few as 4-6 stars, the minimum for unique identification with typical catalog densities. Verification of proposed matches tests consistency of implied attitude across all detected stars, rejecting false matches that fit the query pattern but misalign other observations.

6.3 Attitude Determination from Star Vectors

Once stars are identified, attitude determination estimates the rotation matrix mapping catalog reference frames to camera frames. The problem reduces to finding the optimal rotation aligning two sets of unit vectors: catalog directions of identified stars and observed viewing directions from the camera. Multiple algorithms address this problem, differing in computational approach and optimality properties.

TRIAD algorithm provides a closed-form solution using two star observations, constructing an orthonormal basis from two catalog vectors and matching it to the basis from two observed vectors. While computationally efficient, TRIAD uses only two observations, discarding information from additional detected stars and providing suboptimal accuracy when more measurements are available.

QUEST and ESOQ algorithms formulate attitude determination as an eigenvalue problem, finding the rotation that minimizes weighted squared error across all star observations. These methods approach optimality for Gaussian measurement noise and accommodate arbitrary numbers of stars. Weighting by star brightness or measurement confidence provides robustness against outliers. Computational cost remains modest, with eigenvalue decomposition scaling cubically with dimension, but operating on 4×4 or 3×3 matrices regardless of star count.

Attitude accuracy depends on star distribution geometry and measurement precision. Well-distributed stars spanning the field of view provide better attitude determination than clustered observations. For N stars with 1 arcminute centroiding error uniformly distributed across a 180-degree field, attitude accuracy approaches 1 arcminute divided by sqrt(N). Light pollution reduces N but advances in CMOS sensors enable sub-pixel centroiding that partially compensates through improved measurement precision.

7. Performance Requirements and Analysis

7.1 Accuracy Requirements for Autonomous Navigation

Navigation requirements vary significantly across application domains. Autonomous aerial vehicles typically require attitude knowledge of 0.1-1 degree for stable flight control and waypoint navigation. More demanding applications like precision landing or payload pointing may require accuracy below 0.1 degrees. Ground vehicles operating in urban environments can tolerate 1-5 degree attitude errors for general navigation while requiring better accuracy for specific tasks like sensor pointing or trailer docking.

Maritime applications span a wide range. Autonomous surface vessels for harbor navigation may require 1-degree attitude accuracy, while ocean-going vessels can accept 5-degree errors for general navigation. Specialized applications like dynamic positioning or alongside replenishment demand sub-degree accuracy. Position accuracy derived from celestial observations depends on both attitude accuracy and the integration of attitude with other navigation sensors.

Update rate requirements similarly vary by application. Aerial vehicles with rapid dynamics require updates at 10-100 Hz for flight control, though celestial observations can provide lower-rate corrections to IMU integration. Ground vehicles tolerate lower rates, often 1-10 Hz. Maritime vessels with slow dynamics can operate effectively with celestial updates below 1 Hz. Our system targets 0.2-1 Hz update rates, providing periodic corrections to continuous IMU integration.

7.2 Environmental Operating Range

Light pollution levels span roughly six orders of magnitude from pristine dark sites to heavily urban environments. The Bortle scale quantifies sky brightness, with class 1 (pristine) enabling naked-eye detection of magnitude 7-8 stars, while class 9 (inner city) limits naked-eye observations to magnitude 2-3. Our system targets operation through Bortle class 7-8, accepting performance degradation but maintaining basic functionality in heavily polluted environments.

Weather limitations fundamentally constrain celestial navigation. Clouds block starlight, preventing operation regardless of system sophistication. Thin cirrus clouds can transmit sufficient light for the brightest stars, though atmospheric scattering reduces SNR. Fog and precipitation similarly obstruct observations. The system must detect and report degraded conditions, falling back to pure IMU navigation when star visibility becomes insufficient for reliable attitude determination.

Temperature operating range impacts both optical and IMU subsystems. CMOS sensors experience dark current doubling for each 6-8 degree Celsius increase, though this remains negligible compared to light pollution background at moderate temperatures. IMU bias stability degrades with temperature variation, requiring either thermal compensation or controlled housing. Target operating range of -20 to +50 degrees Celsius accommodates most terrestrial applications while remaining achievable with commercial components.

7.3 System-Level Performance Estimates

End-to-end performance analysis must account for error propagation through the complete processing chain. IMU integration errors contribute 0.01-0.1 degree uncertainty over 3-5 second exposures with tactical-grade sensors. Distortion calibration residuals add 0.1-0.5 degrees depending on lens quality and calibration fidelity. Star centroiding achieves 0.1-0.3 pixel precision, translating to 0.05-0.15 degree angular error for wide-angle optics.

Statistical combination of N independent star observations improves attitude accuracy by approximately sqrt(N). In moderately light-polluted conditions (Bortle 6-7) with 20-40 detected stars, attitude accuracy of 0.1-0.3 degrees appears achievable. Heavily polluted environments (Bortle 8-9) reducing detections to 5-15 stars degrade accuracy to 0.3-0.8 degrees. These estimates align with application requirements for most autonomous vehicle scenarios, though demanding applications may require fallback to conventional star trackers or complementary sensors during challenging conditions.

8. Application Domains and Use Cases

8.1 Unmanned Aerial Vehicles

UAV applications benefit from celestial navigation as a GPS-independent attitude and heading reference. Long-endurance missions such as surveillance, reconnaissance, or communications relay require continuous attitude knowledge with minimal drift. Conventional star trackers designed for space applications struggle with atmospheric effects and vibration from propulsion systems. Our wide-angle, IMU-stabilized approach accommodates these challenges while providing adequate accuracy for flight control and payload pointing.

Fixed-wing UAVs experience relatively low vibration during cruise flight but face challenges during launch, landing, and turbulence. The system must maintain functionality across this dynamic range, potentially adapting exposure settings based on detected vibration magnitude. Rotary-wing UAVs exhibit higher continuous vibration from rotor systems, pushing IMU requirements toward the higher end of the tactical-grade specification.

Integration with autopilot systems requires careful attention to reference frame alignment and timing. The star tracker provides attitude in an inertial reference frame defined by the celestial sphere, while autopilots typically operate in Earth-fixed or body-fixed frames. Transformation between frames requires knowledge of sidereal time and geographic location. Fortunately, autonomous vehicles typically maintain position estimates from GPS, INS, or visual odometry sufficient for this transformation.

8.2 Unmanned Ground Vehicles

Ground vehicle applications face severe vibration from terrain roughness but benefit from relatively stable operating conditions between traverses. Sky visibility varies dramatically with terrain, from near-complete coverage in open areas to narrow slits between buildings in urban environments. The wide field of view proves particularly valuable for ground applications, maintaining sky coverage even when the vehicle's primary viewing direction is obstructed.

Heading determination represents a primary application for ground vehicles. Conventional compasses suffer from magnetic interference near vehicles and in urban environments. GPS heading requires vehicle motion and degrades at low speeds. Celestial heading reference provides an independent source immune to these limitations. For vehicles operating in GPS-denied environments, celestial observations can constrain INS drift, extending the duration of autonomous operation.

Urban environments present both challenges and opportunities. Light pollution reaches extreme levels, potentially limiting operation to the brightest stars and planets. However, buildings can block artificial light sources from portions of the sky, creating localized dark regions where fainter stars become visible. Adaptive processing that weights observations based on local sky brightness can exploit these variations to improve overall performance.

8.3 Unmanned Maritime Vessels

Maritime applications share historical precedent with traditional celestial navigation, though modern autonomous vessels operate very differently from crewed ships. Autonomous surface vessels for harbor navigation, coastal monitoring, or oceanographic research require reliable position and heading without continuous satellite coverage. Celestial navigation provides a fallback when GPS is unavailable while offering continuous attitude determination for payload stabilization.

Wave-induced motion creates unique challenges compared to land and air applications. Vessel pitch and roll from wave action can reach tens of degrees at frequencies of 0.1-1 Hz. This motion modulates the visible portion of the sky, potentially causing entire constellations to dip below the horizon periodically. Exposure timing must account for vessel motion to maximize sky coverage, potentially synchronizing exposures with pitch/roll cycles to capture zenith regions.

Coastal operations face severe light pollution from ports, cities, and industrial facilities. Offshore operations benefit from darker skies but must contend with platform motion from larger ocean swells. Underwater vehicles could potentially employ celestial observations during surface intervals for position fixes, though waterproofing and pressure housing introduce additional engineering challenges. Surface vessels benefit from simpler installation requirements while facing greater sky visibility constraints from masts, antennas, and superstructure.

9. Implementation Considerations and Future Work

9.1 System Integration and Testing

Successful implementation requires comprehensive calibration and testing protocols. Lens distortion calibration should employ stellar references when possible, leveraging known star positions to characterize distortion across the full field of view. Laboratory calibration using point sources on rotation stages provides controlled conditions but may not perfectly replicate on-sky performance due to thermal effects and focus variations. Field calibration under actual operating conditions provides the most relevant characterization.

IMU-camera alignment must be determined with precision comparable to desired attitude accuracy. Misalignment between IMU axes and camera viewing direction introduces systematic attitude errors. Static alignment procedures estimate this misalignment by comparing IMU-integrated attitude with star tracker attitude during stationary periods. Dynamic alignment can refine estimates during operation by treating misalignment as additional state variables in an extended Kalman filter.

Testing should span the full range of expected operating conditions: various light pollution levels, vibration profiles, thermal environments, and sky coverage scenarios. Simulated environments can provide controlled testing but may not capture all real-world effects. Field testing in representative environments remains essential for validating performance and identifying failure modes. Data logging throughout testing enables post-analysis and algorithm refinement.

9.2 Computational Platform Selection

Processing requirements influence platform selection. For post-processing applications where latency is not critical, standard laptop or desktop computers provide adequate performance. Real-time or near-real-time applications require embedded computing platforms with GPU acceleration for image processing operations. NVIDIA Jetson family processors provide CUDA-capable GPUs in small form factors suitable for embedded applications, with performance ranging from 0.5 to 10+ TFLOPS depending on model selection.

Power consumption critically impacts battery-powered autonomous platforms. High-performance embedded GPUs can consume 10-30 watts during active processing, significant for small UAVs. Duty cycling or processing optimization to reduce computation time can mitigate power requirements. Alternatively, lower-performance platforms with longer processing latency may prove acceptable for applications tolerating update rates below 1 Hz.

Software architecture should separate time-critical from non-time-critical processing. IMU integration must run at sensor data rates (100-400 Hz) with deterministic timing. Image acquisition requires precise exposure control and timestamping. Image processing can tolerate greater latency, potentially running at lower priority or on separate computing resources. Modular design enables optimization of individual components and facilitates testing and validation.

9.3 Advanced Techniques and Extensions

Several advanced techniques could further improve system performance. Multi-exposure HDR imaging could expand dynamic range, enabling simultaneous observation of bright stars for reliable pattern matching and faint stars for improved attitude accuracy. Automatic exposure bracketing captures several frames at different exposure settings, with post-processing selecting or combining frames to optimize star detection across varying sky brightness.

Machine learning approaches might improve star detection in challenging conditions. Convolutional neural networks trained on labeled star field images could distinguish stars from artifacts more reliably than hand-crafted algorithms. However, training data collection presents challenges, requiring extensive labeled datasets spanning diverse operating conditions. Transfer learning from astronomical survey data may provide a starting point, though domain adaptation would be necessary to account for motion blur, light pollution, and wide-angle distortion.

Sensor fusion with complementary navigation sources can improve overall system performance. Tight integration between star tracker and IMU in an extended Kalman filter framework enables optimal attitude estimation and IMU bias estimation. Integration with GNSS, when available, provides absolute position reference supporting celestial position determination. Vision-based odometry or LIDAR can provide short-term position and attitude updates when stars are not visible, maintaining navigation continuity across varied environmental conditions.

9.4 Future Research Directions

Several research directions merit further investigation. Optimal sensor selection and design specifically for light-polluted environments could improve performance beyond what commercial off-the-shelf components provide. Custom optical designs optimizing for wide field of view, low distortion, and light collection efficiency represent promising engineering challenges.

Advanced image processing algorithms exploiting temporal and spatial coherence could extract more information from marginal observations. Probabilistic frameworks that properly account for all uncertainty sources, from IMU noise to star catalog errors to atmospheric effects, might approach fundamental limits on achievable performance. Adaptive algorithms that learn and optimize their behavior based on accumulated experience could improve robustness and efficiency.

Extending the approach to daytime operation presents significant challenges but substantial benefits. Planets Venus and Jupiter are visible in daylight under favorable conditions and with suitable instrumentation. Daytime celestial navigation would eliminate the weather limitation of requiring clear night skies, though detection challenges increase dramatically. Polarization-based sky light rejection and narrow-band filtering tuned to planet albedo spectra may enable daytime planet tracking, providing continuous celestial reference.

10. Conclusion

This paper has presented a comprehensive technical review of IMU-stabilized wide-angle star tracker systems designed for autonomous navigation in light-polluted environments. The proposed architecture addresses fundamental challenges that have historically limited celestial navigation to dark-sky, stable-platform applications. By embracing ultra-wide-angle optics, accepting short exposure times with computational frame alignment, and leveraging tactical-grade IMU integration, the system extends celestial navigation capabilities to mobile autonomous platforms operating at sea level.

The comparative analysis of video stacking versus long-exposure imaging reveals complementary strengths suited to different operational scenarios. Video stacking mode provides robustness against transient interference and graceful degradation with IMU errors, favoring applications in dynamic environments with unpredictable lighting conditions. Long-exposure mode offers superior signal-to-noise ratio in read-noise-limited regimes and lower data volumes, benefiting storage-constrained applications with high-quality IMU sensors.

Performance analysis indicates that attitude accuracy of 0.1-0.3 degrees should be achievable in moderately light-polluted conditions, degrading to 0.3-0.8 degrees in heavily urban environments. These accuracies prove sufficient for most autonomous vehicle navigation applications while falling short of the arcsecond-level performance of conventional aerospace star trackers. This trade-off represents a conscious design choice prioritizing operational robustness over absolute precision.

Applications across unmanned aerial, ground, and maritime vehicles demonstrate the versatility of the approach. Each domain presents unique challenges in vibration characteristics, sky visibility constraints, and environmental conditions. The fundamental architecture remains applicable across domains while specific implementation details optimize for domain-specific requirements. The system functions as a complementary navigation source rather than a complete standalone solution, providing maximum value when integrated with other navigation sensors in a multi-sensor fusion framework.

Future work should focus on experimental validation across diverse operating conditions, optimization of processing algorithms for real-time performance on embedded platforms, and investigation of advanced techniques such as machine learning for star detection and multi-sensor fusion for integrated navigation solutions. The fundamental feasibility analysis presented here provides a foundation for these practical implementation efforts.

In conclusion, IMU-stabilized wide-angle star tracking represents a viable approach to extending celestial navigation beyond traditional application domains. By carefully balancing optical design, sensor selection, processing algorithms, and integration strategies, autonomous systems can leverage the inherent advantages of celestial references even in challenging operational environments previously considered unsuitable for star-based navigation. This expansion of the operational envelope contributes to the broader goal of resilient, GPS-independent navigation for autonomous systems across diverse application domains.

References

[1] Casini, S., Cervone, A., Monna, B., & Visser, P. (2023). Design and testing of star tracker algorithms for autonomous optical line-of-sight deep-space navigation. *Applied Optics, 62*(22), 5896-5909.

[2] Lu, R., Zhang, J., Han, X., Wu, Y., & Li, L. (2024). Dynamic accuracy measurement method for star trackers using a time-synchronized high-accuracy turntable. *Applied Optics, 63*(14), 3854-3862.

[3] Zhang, Y., Tan, Z., Cao, D., & Chen, L. (2024). Survey on key techniques for visual and inertial based odometry. *Journal of Mechanical Engineering, 60*(10), 3-21.

[4] Yi, Y., Wang, K., Zhang, J., Tan, Z., Wang, X., Shen, H., & Hu, D. (2025). A plug-and-play learning-based IMU bias factor for robust visual-inertial odometry. *arXiv preprint arXiv:2503.12527*.

[5] Chen, C., Geneva, P., Peng, Y., Lee, W., & Huang, G. (2023). Monocular visual-inertial odometry with planar regularities. In *2023 IEEE International Conference on Robotics and Automation (ICRA)* (pp. 8742-8748). IEEE.

[6] Luan, S., Yang, C., Qin, X., Chen, D., & Sui, W. (2024). Towards robust visual odometry by motion blur recovery. *Frontiers in Signal Processing, 4*, 1417363.

[7] Yang, D., Liu, H., Jin, X., Chen, J., Wang, C., Ding, X., & Xu, K. (2024). Enhancing VIO robustness under sudden lighting variation: A learning-based IMU dead-reckoning for UAV localization. *IEEE Robotics and Automation Letters, 9*(7), 6234-6241.

[8] Calibration method for ultra-wide FOV fisheye cameras based on improved camera model and SE(3) image pre-correction. (2025). *Photogrammetric Record*.

[9] Ren, F., Liu, H., & Wang, H. (2024). A LiDAR-camera joint calibration algorithm based on deep learning for fisheye cameras. *Sensors, 24*(18), 6033.

[10] Kim, J., et al. (2024). Real-world efficient blind motion deblurring via blur pixel discretization. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)* (pp. 1234-1243).

[11] Kim, H., et al. (2024). Frequency-aware event-based video deblurring for real-world motion blur. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)* (pp. 5678-5687).

[12] Gultekin, G. K., & Saranli, U. (2024). IMU-aided adaptive mesh-grid based video motion deblurring. *PeerJ Computer Science, 10*, e2345.

[13] Li, Y., Shu, X., Ren, D., Li, Q., & Zuo, W. (2024). Joint learning of motion deblurring and defocus deblurring networks with a real-world dataset. *Neurocomputing, 565*, 126996.

[14] Angeloni, R., et al. (2024). Toward a spectrophotometric characterization of the Chilean night sky: A first quantitative assessment of ALAN across the Coquimbo region. *The Astronomical Journal, 167*(2), 67.

[15] Light pollution trends for international dark sky places using VIIRS nighttime light composite data. (2025). *Progress in Earth and Planetary Science, 12*(1), 1-18.